

Jun-Kun Chen

Ph.D. Candidate in Computer Science | Expected Dec. 2026

immortalco233@gmail.com | junkun3@illinois.edu | +1 872-772-0116 | [Website](#) | [Scholar](#) | [LinkedIn](#)

Legal name: Junkun Chen | Publications: Jun-Kun Chen / J.-K. Chen

Research Focus

Research in **diffusion-based controllable generation**: multimodal conditioning, temporally consistent image/video synthesis, and visual world models. Current work connects language reasoning with **visual outcome prediction** through multimodal post-training.

Education

University of Illinois Urbana-Champaign

Ph.D. Candidate in Computer Science; Advisor: Prof. **Yu-Xiong Wang**

Master of Computer Science, Dec. 2023 (earned during Ph.D. study)

Tsinghua University

B.Eng. in Computer Science, Yao Class (IIIS); Major GPA: 3.94/4.0

2021 – Dec. 2026 (expected)

Urbana, IL

2017 – 2021

Beijing, China

Selected Publications

1. Virtual Fitting Room: Generating Arbitrarily Long Videos of Virtual Try-On from a Single Image [↗](#) **NeurIPS 2025**
Jun-Kun Chen, Aayush Bansal, Minh Phuoc Vo, Yu-Xiong Wang [Project](#) | [Early preview \(arXiv, Sep. 2025\)](#)
2. ProEdit: Simple Progression is All You Need for High-Quality 3D Scene Editing [↗](#) **NeurIPS 2024**
Jun-Kun Chen, Yu-Xiong Wang [Project](#) | [arXiv \(Nov. 2024\)](#)
3. ConsistDreamer: 3D-Consistent 2D Diffusion for High-Fidelity Scene Editing [↗](#) **CVPR 2024**
Jun-Kun Chen, Samuel Rota Bulò, Norman Müller, Lorenzo Porzi, Peter Kotschieder, Yu-Xiong Wang [Project](#) | [arXiv \(Jun. 2024\)](#)
4. Instruct 4D-to-4D: Editing 4D Scenes as Pseudo-3D Scenes Using 2D Diffusion [↗](#) **CVPR 2024**
Linzhan Mou*, **Jun-Kun Chen***, Yu-Xiong Wang [Project](#) | [arXiv \(Jun. 2024\)](#)
5. HDEdit: Editing Videos and 3D Scenes with Video Diffusion Through Hierarchical Task Decomposition [↗](#) **3DV 2026**
Yanming Zhang*, **Jun-Kun Chen***, Jipeng Lyu, Yu-Xiong Wang [Project](#) | [Earlier version \(arXiv, Mar. 2025\)](#)
6. Dress&Dance: Dress up and Dance as You Like It [↗](#) **WACV 2027**, under review
Jun-Kun Chen, Aayush Bansal, Minh Phuoc Vo, Yu-Xiong Wang [Project](#) | [Early preview \(arXiv, Aug. 2025\)](#)

*Equal contribution. Earlier publications also appear as Junkun Chen / J. Chen.

Industry Experience

Meta Reality Labs

Research Scientist Intern (2025 continued part-time Aug.–Dec.)

May–Dec. 2025; May–Aug. 2023

Bay Area, CA, U.S. / Zurich, Switzerland

- **2025**: Designed parallel training workflows for 3D/video generation at **128–256-GPU scale** across multi-node clusters.
- **2023**: Developed multi-view conditioning, structured noise, and self-supervised consistency training for diffusion editing; published as *ConsistDreamer* at CVPR 2024.

SpreeAI

AI Research Intern

May 2024 – May 2025

Remote, U.S.

- **Dress&Dance**: designed **multimodal conditioning** and progressive image/video training for controllable 5-second try-on videos at **1152×720, 24 FPS**.
- **Virtual Fitting Room** (NeurIPS 2025): extended this work to **minute-scale, identity-consistent generation** with anchored autoregression; evaluated up to 90 seconds using short-clip training.

Current Research: Generative World Models

Research direction: connecting language reasoning and visual generation through training grounded in predicted outcomes.

Computer Use — Action-Conditioned Image-Diffusion World Model

- **Post-trained image diffusion** on SenseNova MoT for next-screen prediction conditioned on the current screenshot and action.
- Designed **action-consequence and content-preservation supervision**; incorporated interaction traces and failure cases into Reasoner training.

Physics — Video-Diffusion World Model for Physical Events

- Adapted **Cosmos3 video diffusion** for physical-event prediction from observed RGB history.
- Implemented **online teacher distillation**, training the language Reasoner and diffusion Generator from video prediction errors.

Robotics — Motion-Conditioned Video Diffusion World Model

- Built an initial synthetic **video-diffusion prototype** for robot-object interaction prediction from a scene and prescribed motion.
- Designed **visual-proxy and paired-correspondence conditioning** to represent motion across robot-world setups.

Technical Skills & Honors

Tools & Methods: PyTorch; Python; C++; image/video diffusion; diffusion Transformers (DiT); LoRA; vision-language models; multimodal conditioning; training-data construction; benchmark design; quantitative and human evaluation; diffusion post-training.

Competition Highlights: ICPC EC-Final Gold (4th); ICPC Qingdao Gold (2nd); CCPC Harbin Gold; NOI Gold.

Selected Research Contributions

Controllable Try-On to Long-Horizon Video — Dress&Dance and VFR

- **Dress&Dance**: designed attention-based **multimodal conditioning** and progressive image/video training; achieved **20–59% lower FVD** than a Kling 1.6 try-on-and-animation pipeline on Internet and captured benchmarks (full manuscript, Table C.4).
- Constructed garment/user/video training triplets using **synthetic appearance conditions with real-video targets**; staged the conditioning and resolution curricula to teach new controls while retaining pretrained visual priors.
- **Virtual Fitting Room**: designed **anchor/prefix conditioning** and immediate refinement for long-video generation from short training clips; scored higher on VBench subject/background consistency and motion smoothness than D&D+FramePack and D&D+Kling 2.0 pipelines on 30–50 s hard-motion clips.

General-Purpose Diffusion Editing — HDEdit and ProEdit

- **HDEdit**: designed **hierarchical task decomposition** that separates instruction fulfillment from content preservation, combining LLM-based scheduling with interpolation-based refinement for training-free video and scene editing.
- **ProEdit**: developed tuning-free preservation control and a **progressive editing workflow** with preview and edit-strength control, enabling high-fidelity geometry and appearance editing.

Multi-View and Temporal Consistency — ConsistDreamer and Instruct 4D-to-4D

- **ConsistDreamer**: created context-rich multi-view conditioning, **3D-consistent structured noise**, and self-supervised warping consistency for high-fidelity diffusion editing across views.
- **Instruct 4D-to-4D**: extended 2D diffusion to long-term dynamic scenes through **anchor-aware attention** and optical-flow-guided sliding windows.

Geometric Visual Representations — NeuralEditor and PointTree

- Built point-based radiance-field and relaxed k-d-tree representations for editable scene geometry and **transformation-robust point-cloud encoding**.

Additional Publications

7. SceneCraft: Layout-Guided 3D Scene Generation  NeurIPS 2024
Xiuyu Yang*, Yunze Man*, **Jun-Kun Chen**, Yu-Xiong Wang Project | [arXiv \(Oct. 2024\)](#)
8. Contrastive Learning Relies More on Spatial Inductive Bias Than Supervised Learning: An Empirical Study  ICCV 2023
Yuanyi Zhong*, Haoran Tang*, **Jun-Kun Chen***, Yu-Xiong Wang
9. NeuralEditor: Editing Neural Radiance Fields via Manipulating Point Clouds  CVPR 2023
Jun-Kun Chen*, Jipeng Lyu*, Yu-Xiong Wang
10. PointTree: Transformation-Robust Point Cloud Encoder with Relaxed K-D Trees  ECCV 2022
Jun-Kun Chen, Yu-Xiong Wang
11. Is Self-Supervised Contrastive Learning More Robust Than Supervised Learning?  ICML 2022 Workshop
Yuanyi Zhong*, Haoran Tang*, **Jun-Kun Chen**, Jian Peng, Yu-Xiong Wang
12. TorchDrug: A Powerful and Flexible Machine Learning Platform for Drug Discovery  arXiv preprint, 2022
Zhaocheng Zhu, Chence Shi, Zuobai Zhang, ..., **Jun-Kun Chen**, ..., Jian Tang
13. RNNLogic: Learning Logic Rules for Reasoning on Knowledge Graphs  ICLR 2021
Meng Qu*, **Jun-Kun Chen***, Louis-Pascal Xhonneux, Yoshua Bengio, Jian Tang
14. Pre-training Assisted Scene-aware Dialogue Generation  DSTC8 Workshop at AACL 2020
Jun-Kun Chen, Huang He, Siqi Bao, Fan Wang

Earlier Experience

- | | |
|---|------------------------------|
| Mila – Quebec AI Institute | Feb. – Aug. 2020 |
| Research Intern; knowledge-graph reasoning with Prof. Jian Tang | Montreal, Canada |
| Baidu, Natural Language Processing Department | July 2019 – Jan. 2020 |
| Research / Development Intern; scene-aware dialogue generation (DSTC8 top-3 system) | Beijing, China |

Competitions, Leadership & Coursework

Programming Competitions: ICPC EC-Final 2018, Gold Medal (4th); ICPC Qingdao Regional 2018, Gold Medal (2nd); CCPC Harbin Regional 2017, Gold Medal; NOI 2016, Gold Medal.

Additional Distinctions: CTSC 2017 (7th); CTSC 2016 Gold (8th); NOI Winter Camp 2016 Gold (3rd).

Leadership & Problem Setting: Problem setter for NOI, CTSC, ICPC, and CCPC events; Vice Chair, Tsinghua Student Association of Algorithms and Competitions (2018–2019).

Selected Coursework (Generative Models): Deep Generative Models (A+); Generative Models for Biomedicine and Life Sciences (A).

Vision & 3D: 3D Vision (A+); Deep Learning for Computer Vision (A); Computer Vision (A); Computational Photography (A+).

Learning: Meta-Learning (A+); Machine Learning (A).